



The AI Landscape: Exploring the Latest Technologies

Joseph M. Reilly, PhD

Assistant Teaching Professor

College of Professional Studies

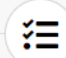


Please download the demo materials!

Session 2: The AI Landscape: Exploring
the Latest Technologies




Materials & Tasks



Materials & Tasks

- Read [Democratising AI: Multiple Meanings, Goals, and Methods](#) with a focus on Section 2: Democratisation of AI Use
- Watch [Andrew Ng on The Near Future of AI](#)



- Download a container to run the llm platform: <https://www.docker.com/products/docker-desktop/>
- Download a llm platform: <https://ollama.com/Links to an external site.>
- Open your terminal and type the following command: `docker run -d -p 3000:8080 --add host=host.docker.internal:host-gateway -v open-webui:/app/backend/data --name open-webui --restart always ghcr.io/open-webui/open-webui:main`
- Then press 'Enter'



Agenda for First Portion

1

The Current AI Market (5 mins)

2

Explainable AI and Open Source (10 mins)

3

Low-code and No-code AI (10 mins)

4

Example: Google Vertex AI (5 mins)



The Current AI Market

Where we're at and where we're going (maybe)




AI Market Expansion

2024
\$638b


2034
\$3,680b

- North America largest market share, but other regions catching up
- Heavy investment by tech giants, demand from end-use verticals
- Favorable government initiatives and little regulation
- Lack of skilled professionals to leverage AI
 - Augmenting the workforce, not automating it away

Growth Driven by Increased Access



The Trough of Disillusionment





Explainable AI and Open Source

Explain the behavior and see how we got there

The Danger of the Black Box

Data → Prediction

- If you can't explain why a model arrived at a specific decision, how can you defend it?
- Big issue in deep learning
- How can you trust it?
 - Mortgage lending
 - Medical diagnoses
 - Predictive policing
 - Hiring and firing





Explainable AI (XAI)

Interpretability

- Observe the inner workings
- Cause and effect of each parameter and input
- Very difficult for complex, high performance models

Explainability

- Explain the behavior in human terms beyond weights and features
- Indirect, model-agnostic ways to find meaning from results
- Communicate to end users

Must accuracy be sacrificed for interpretability?
Not necessarily.



Open-Source AI

Harder to define than you'd think!

Current draft from the Open Source Initiative

Free software means the freedom to...

- run the program as you wish
- study how the program works and change it as you wish
- redistribute copies to help others
- distribute copies of your modified versions to others

Complications with AI:

- Free access to the training data?
- Use-based licenses?
- Code related to preprocessing?



Benefits of Open Source

Transparency

Understand how it works
and mitigate biases

Customization

Fine-tune models on
sensitive data

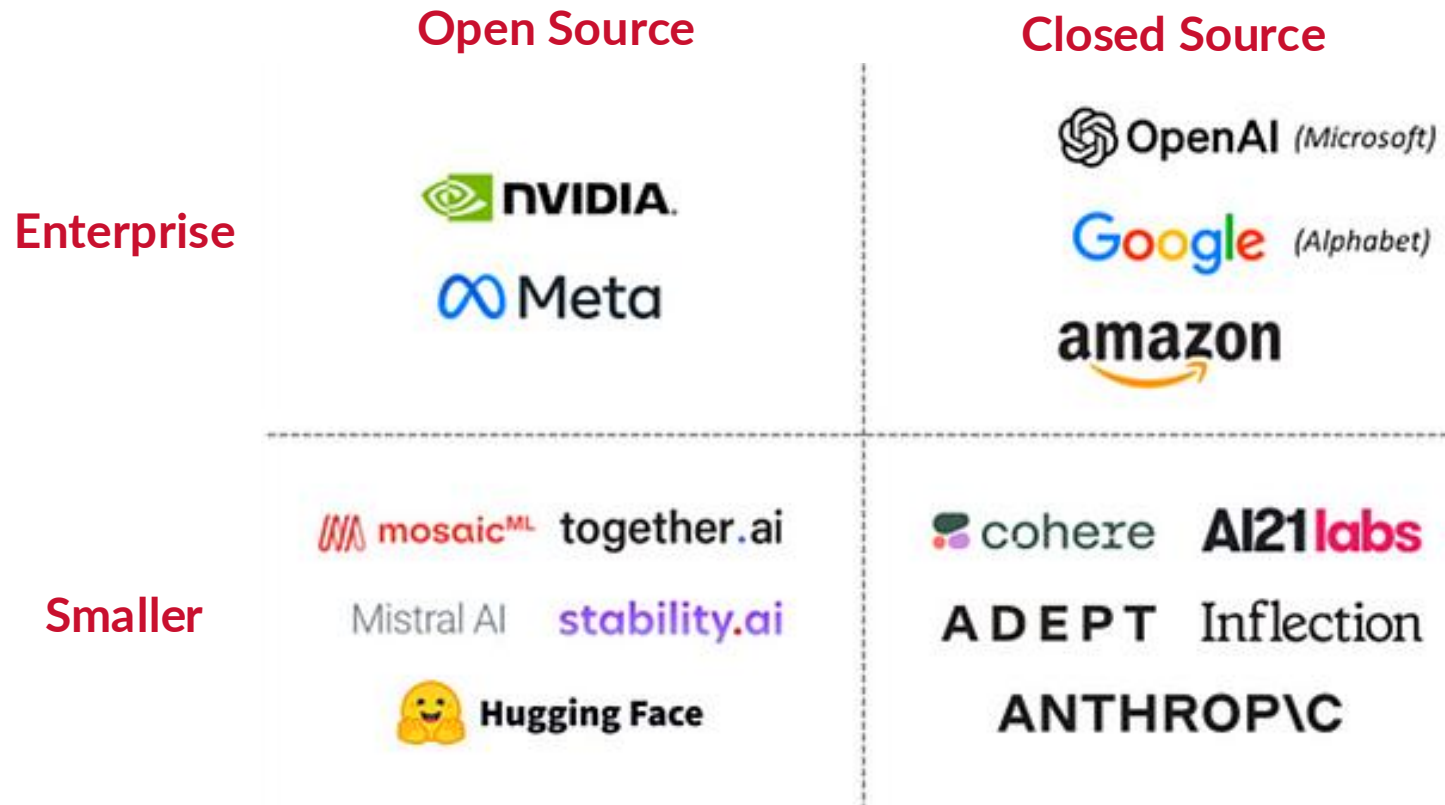
Collaboration

Understand how it works
and mitigate biases

Lower Barriers to Use

Efficiency and affordability
to experiment and deploy

Open-Source AI Examples





Low-code and No-code AI

Democratizing AI and empowering new users



Emergence of Low-Code Options

- Visual approach to development to reduce traditional programming
 - Traces back to late 80's / early 90's with Visual Basic and web development
- Helpful for experienced devs as well as novices
- Connecting with existing ML/AI tools

20% growth
from '22 to '23

80% of user base
outside of IT dept

70% of new apps
will use low code by '25



Advantages of Enterprise Low-Code

Agility

Less Tech Debt

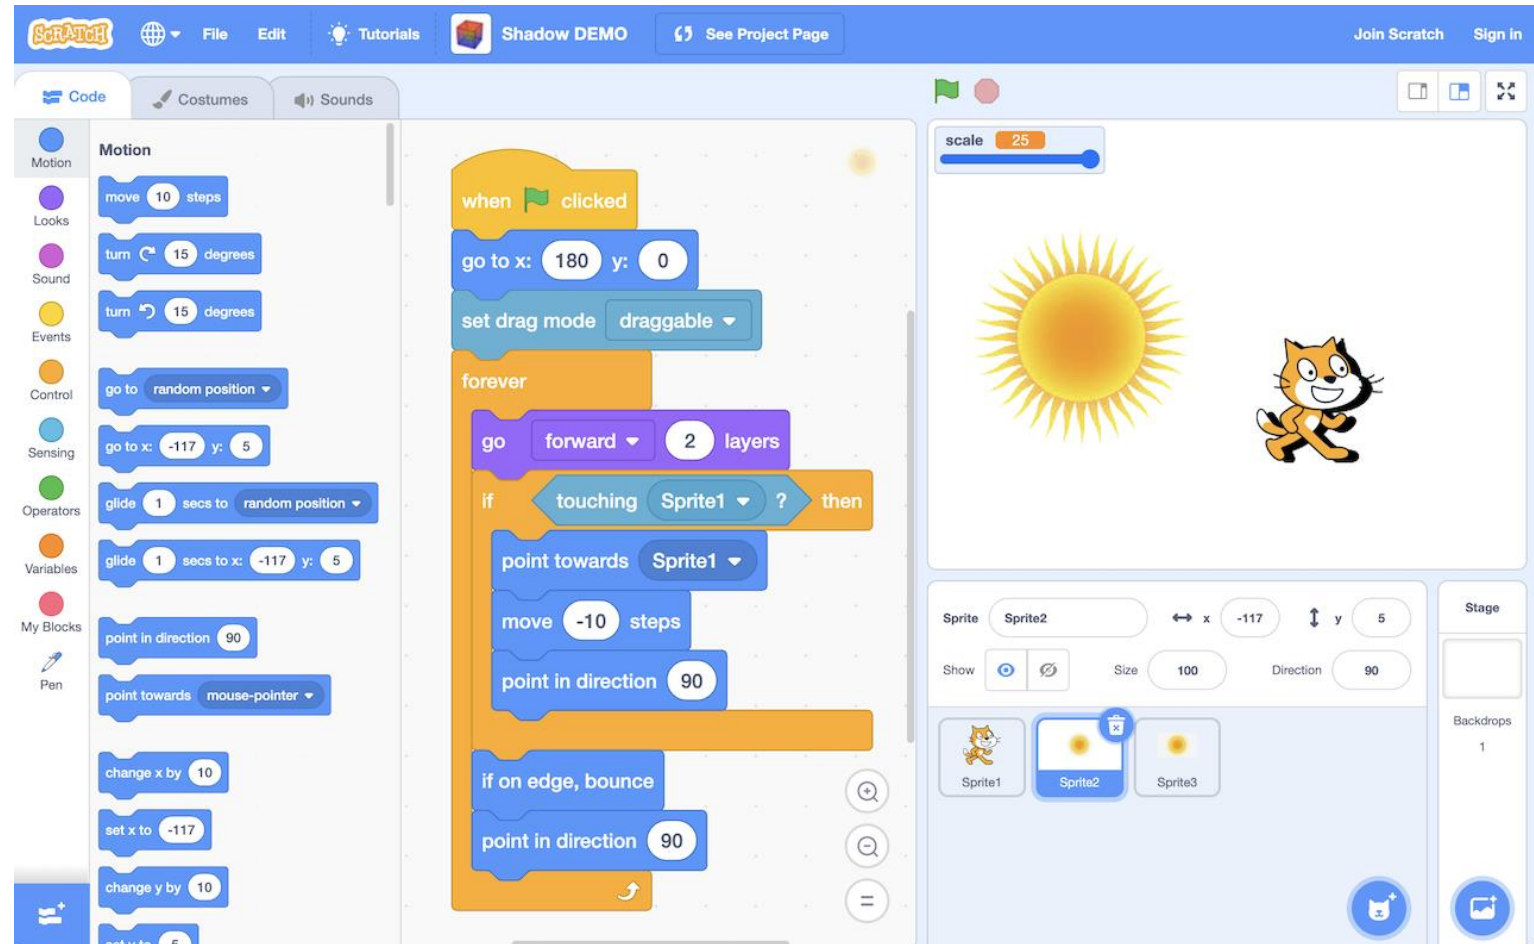
Lower Costs

Speeds Up Dev

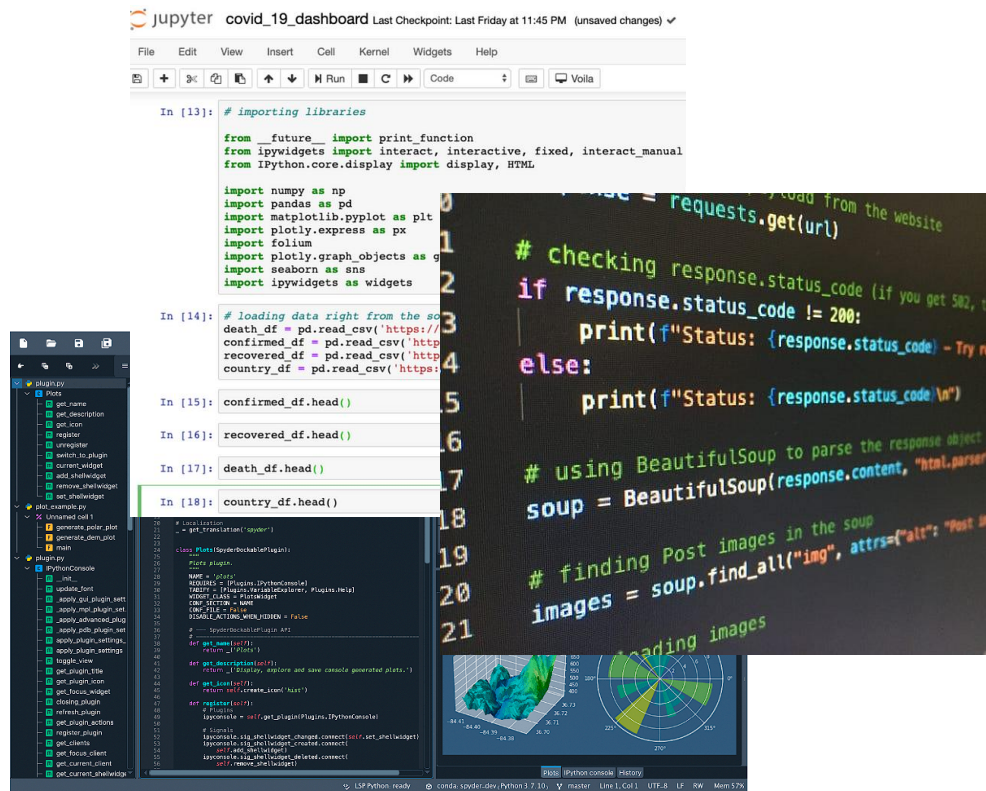
Citizen Developer

Levels of
Abstraction

Scratch!



Evolution of AI Development



The screenshot shows a Jupyter Notebook titled 'covid_19_dashboard'. The code in the cells includes:

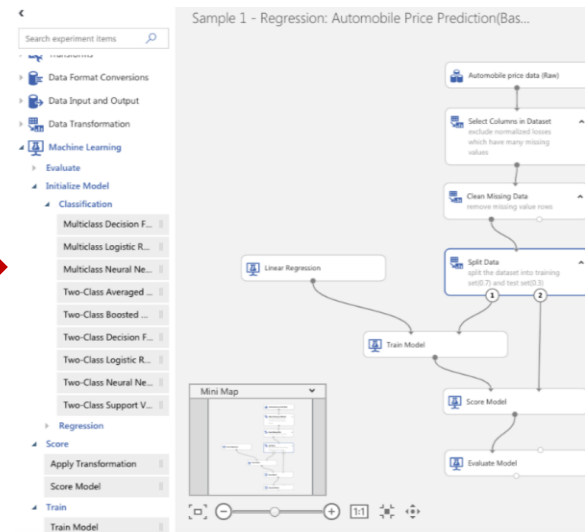
```
In [13]: # importing libraries
from __future__ import print_function
from ipywidgets import interact, interactive, fixed, interact_manual
from IPython.core.display import display, HTML

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import folium
import plotly.graph_objects as go
import seaborn as sns
import ipywidgets as widgets

In [14]: # loading data right from the web
death_df = pd.read_csv('https://...')
confirmed_df = pd.read_csv('http...')
recovered_df = pd.read_csv('http...')
country_df = pd.read_csv('https...')

In [15]: confirmed_df.head()
In [16]: recovered_df.head()
In [17]: death_df.head()
In [18]: country_df.head()
```

Below the code cells, there are two visualizations: a 3D surface plot on the left and a circular heatmap on the right. The bottom status bar indicates the environment is 'conda-tour-de-voy Python 3.7.10'.



Tune hyperparameters of a Random Forest Classifier on my training data, then report the precision and recall on the test set. Explain the results concisely.


Representative Sample of Some Players




Amazon SageMaker



Azure Machine Learning



DataRobot



IBM Watson



Example: Google Vertex AI

See how this works in practice at enterprise scale



Important Caveats!

I do not work for Google.

They aren't paying me to present this.

I don't currently use Vertex AI but I have used GCP.

It all starts with data

Dataset name *
example_dataset

Can use up to 124 characters.

Select a data type and objective


First select the type of data your dataset will contain. Then select an objective, which is the outcome that you want to achieve with the trained model. [Learn more](#)

IMAGE


TABULAR

TEXT


VIDEO



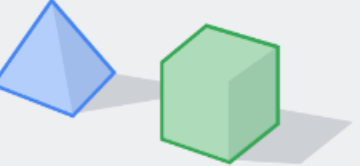
☐ **Single-label classification**
Predict a single label for an image.



☒ **Multi-label classification**
Predict one or more labels for an image



☐ **Object detection**
Predict objects in an image with a label and location (bounding box)



☐ **Segmentation**
Predict per-pixel areas of an image with a label

Region
us-central1 (Iowa)

It all starts with data

IMPORT

BROWSE

ANALYZE

LINEAGE

Labels

All3,667

Labeled3,667

Unlabeled0

Filter

+

daisy633

dandelion898

roses641

sunflowers697

tulips799


ADD NEW LABEL

Images


Filter

Enter label or property name


Select all




roses




daisy




tulips



daisy



sunflowers



dandelion

Related resources

Training jobs and models

Use this dataset and annotation set to train a new machine learning model with AutoML or custom code

TRAIN NEW MODEL

Labeling tasks

If your data still needs to be labeled, create a labeling task to have others label it for you

Vertex AI Data Labeling Service (requesting human labelers) is deprecated and will no longer be available on Google Cloud after July 1, 2024. [Learn more](#)

CREATE LABELING TASK

AI IN ACTION
Building Your Essential AI Toolkit

23

Model training on the cloud

Train new model

- ✓ Training method
- 2 Model details
- 3 Training options
- 4 Compute and pricing

START TRAINING CANCEL

- ☒ Train new model
Creates a new model group and assigns the trained model as version 1
- ☐ Train new version
Trains model as a version of an existing model

Name *
example_model

Description
let's see how this works

Data split

- ☒ Random ☐ Manual (Advanced)

Your dataset will be automatically randomized and split into training, validation, and test sets using the following ratios. [Learn more](#)

Training 80 % Validation 10 % Test 10 %

- Training: 80%
- Validation: 10%
- Test: 10%






Name	ID	Status	Job type	Model type	Duration ?
example_model	5690383994168803328	Training	Training pipeline	Image classification (Multi-label)	1 hr 28 min

Model evaluation

Evaluation details


Confidence threshold   0.5

All labels


Average precision 	0.998
Precision 	98.6%
Recall 	96.5%
Created	Sep 16, 2024, 1:26:47 PM
Total images	3,667
Training images	2,933
Validation images	367
Test images	367

To evaluate your model, set the confidence threshold to see how precision and recall are affected. The best confidence threshold depends on your use case. Read some [example scenarios](#) to learn how evaluation metrics can be used.

Precision-recall curve




Precision-recall by threshold



True label	Predicted label				
	roses	dandelion	daisy	tulips	sunflowers
roses	97%	0%	0%	3%	0%
dandelion	0%	99%	1%	0%	0%
daisy	3%	0%	95%	0%	2%
tulips	1%	0%	0%	99%	0%
sunflowers	1%	0%	0%	0%	99%


Model deployment



Model deployment

Test your model **PREVIEW**

Test your model interface showing a photo of an elephant and a list of predicted labels with confidence scores.



Filter	Filter labels
roses	0.210
dandelion	0.129
daisy	0.252
tulips	0.395
sunflowers	0.131



AI IN ACTION

○ Building Your Essential AI Toolkit

Thank you!